

Machine learning under physical constraints

Introduction to RNN

Sixin Zhang
(sixin.zhang@toulouse-inp.fr)

Outline

Recurrent Neural Networks (RNN)

Training strategies of RNN

Outline

Recurrent Neural Networks (RNN)

Training strategies of RNN

RNN for sequence processing

- ▶ Many data such as time-series, language, speech, genomics can be represented in a form of sequence: x_1, x_2, \dots .
- ▶ How to process such data of various type and length?
- ▶ Two basic ideas: **recurrent and convolutional**.
- ▶ **Goal of this lecture:** define what is RNN, and to construct, train and use it.
- ▶ Reference: Ian Goodfellow, Yoshua Bengio, Aaron Courville. Deep Learning. MIT Press, 2016

What is Recurrent?

- ▶ View 1: Output of a dynamical system is fed back to some of its inputs, e.g. time-delayed system.

$$\frac{d}{dt}x(t) = f(t, x(t), x(t - \tau)), \quad \tau > 0$$

- ▶ View 2: Finite-state automata (or Turing machine)
 - ▶ Change from one state to another in response to some inputs
 - ▶ Computers are recurrent !

RNN: dynamical system view

- ▶ Assume state of a dynamical system at time t is h_t
- ▶ Classical form (θ is a parameter)

$$h_t = f(h_{t-1}; \theta)$$

- ▶ Input (external-signal) dependent form

$$h_t = f(h_{t-1}, x_t; \theta)$$

- ▶ h_t can be interpreted as hidden states in NN.

RNN: Automata view

- Automata: an abstract machine that can be in exactly one of a finite number of states at any given time




Figure: Representation of an acceptor; this example shows one that determines whether a binary number has an even number of 0s, where S_1 is an accepting state and S_2 is a non accepting state. See https://en.wikipedia.org/wiki/Finite-state_machine

Unrolling a dynamical system

- ▶ h_t depends on x_t, x_{t-1}, \dots, x_1 and h_0 .
- ▶ To see this, unroll recursively the hidden states:

$$h_t = f(h_{t-1}, x_t; \theta) = f(f(h_{t-2}, x_{t-1}; \theta), x_t; \theta)$$




Figure: Unrolled computational graph

Definition of RNN

- ▶ RNN: a family of NN constructed from the idea of unrolling.
- ▶ Several examples:




Figure: One hidden-layer RNN: allow the network's hidden units to see its (own) previous output

Example 1

- ▶ A concrete case of one hidden-layer RNN:

$$a_t = Ux_t + b + Wh_{t-1}$$

$$h_t = \tanh(a_t)$$

$$o_t = c + Vh_t$$

$$\hat{y}_t = \text{softmax}(o_t)$$

- ▶ Input-to-hidden parameters: U, b
- ▶ Hidden-to-hidden parameters: W, b
- ▶ Hidden-to-output parameters: V, c
- ▶ A loss L_t is then computed based on y_t and \hat{y}_t .

Example 2




Figure: Allow the network's hidden units to see the previous output

Example 3




Figure: Feedback only come from the final output

Example 3

- ▶ A concrete case of TP
 - ▶ Goal: learn an optimal output o_T close to the target y_T .
 - ▶ Initial input $x_0 \sim p = \mathcal{N}(\mu, \sigma^2)$
 - ▶ $h_0 = x_0$
 - ▶ $h_t = h_{t-1} + \theta$
 - ▶ $o_T = h_T$
 - ▶ $y_T = \mu$
 - ▶ $L_T = (y_T - o_T)^2$
 - ▶ Optimization problem: $\min_{\theta \in \mathbb{R}} \mathbb{E}_{x_0 \sim p}(L_T)$

Elman RNN

- ▶ Finding structure in time (Elman 1990)
- ▶ Represent **time** by the effect it has on processing.
- ▶ Applications: sequential prediction, language understanding.




Figure: Save effects of time in a context state, aka memory

Example of Elman RNN: Data assimilation networks




Figure: Unroll Elman RNN over time

- ▶ Relation with Elman RNN and Data assimilation networks
 - ▶ Hidden h_t : analysis posterior distribution
 - ▶ Context c_t : prediction posterior distribution
 - ▶ Input x_t : observed state in ODS

Advanced RNN

- ▶ Next-character generation (basis of **Chat-GPT**): **demo**
- ▶ Multiple hidden-layer (deep) RNN: audio **source separation** (demo)
- ▶ Bi-directional RNN: language model (read a sequence in 2 directions)
- ▶ Auto-encoder RNN: sequence to sequence language **translation**
- ▶ GRU, LSTM: long-dependency in signal (**RNNnoise**: demo)

Outline

Recurrent Neural Networks (RNN)

Training strategies of RNN

Back-propagation over time (BPTT)

- ▶ To train a RNN, one often uses the gradient of the parameters for efficient optimization. BPTT is a way to compute such gradients.
- ▶ It is based on the same idea of back-propagation in neural networks, but it is subtle due to the shared parameters across time.
- ▶ Nowadays, one can use automatic differentiation to do this. But one still needs to understand it to go further.

Example 1

- ▶ The concrete case

$$a_t = Ux_t + b + Wh_{t-1}$$

$$h_t = \tanh(a_t)$$

$$o_t = c + Vh_t$$

$$\hat{y}_t = \text{softmax}(o_t)$$

- ▶ What is the gradient of the softmax layer, $\hat{y}_t = \text{softmax}(o_t)$?
- ▶ A loss L_t is computed based on y_t and \hat{y}_t .
- ▶ How to compute the total loss $L = \sum_t L_t$ with respect to all the parameters?

Example 1: BPTT by chain rule

- ▶ Total loss $L = \sum_t L_t$

$$\frac{\partial L}{\partial L_t} = 1$$

- ▶ From L_t to o_t

$$\nabla_{o_t} L = \left(\frac{\partial L}{\partial o_t} \right)^T = \left(\frac{\partial L}{\partial L_t} \frac{\partial L_t}{\partial o_t} \right)^T = \left(\frac{\partial L_t}{\partial o_t} \right)^T$$

- ▶ From o_t to h_t : output $o_t = c + Vh_t$

$$\nabla_{h_t} L = V^T \nabla_{o_t} L$$

Example 1: BPTT by chain rule

- ▶ Compute gradient of h_t from h_{t+1} .
- ▶ Hidden states:

$$h_{t+1} = \tanh(Ux_{t+1} + b + Wh_t)$$

- ▶ Recursive relation

$$\nabla_{h_t} L = \left(\frac{\partial h_{t+1}}{\partial h_t} \right)^T \nabla_{h_{t+1}} L + \left(\frac{\partial o_t}{\partial h_t} \right)^T \nabla_{o_t} L$$

Example 1: BPTT by chain rule

- ▶ How about the parameters, e.g. (U, b, W) ?

$$h_t = \tanh(Ux_t + b + Wh_{t-1})$$

- ▶ Idea: **accumulate all the gradients over t .**
- ▶ Write b as b_t to be clear of the partial derivatives:

$$\nabla_b L = \sum_t \left(\frac{\partial h_t}{\partial b_t} \right)^T \nabla_{h_t} L$$

where $h_t = \tanh(Ux_t + b_t + Wh_{t-1})$.

Deterministic (BPTT) and online (truncated BPTT)

- ▶ Initialize $\theta = \theta^{(0)}$
- ▶ Deterministic optimizer (BPTT) at each iteration k :

$$\theta^{(k+1)} = \theta^{(k)} - \eta_k \nabla_{\theta} L(\theta^{(k)})$$

- ▶ Online optimizer (truncated BPTT) at iteration k :

$$\theta^{(k+1)} = \theta^{(k)} - \eta_k \tilde{\nabla}_{\theta} L_{k+1}(\theta^{(k)})$$

Example 1: Truncated BPTT

- ▶ The cost (both CPU and memory) to compute $\nabla_b L$ is $O(T)$ due to the summation over $t \leq T$. This is prohibitive when T is very large.
- ▶ Truncated BPTT reduces this cost by focusing on the impact of the “current” parameter b_t on the current loss L_t .
- ▶ The truncated gradient of b at time t is

$$\tilde{\nabla}_b L_t = \left(\frac{\partial h_t}{\partial b_t} \right)^T \nabla_{h_t} L_t, \quad \nabla_{h_t} L_t = \left(\frac{\partial o_t}{\partial h_t} \right)^T \nabla_{o_t} L_t$$

- ▶ The cost to compute $\tilde{\nabla}_b L_t$ is $O(1)$.
- ▶ Discussion of p -truncated BPTT ($p = 1, 2, \dots$) in the paper: Corentin Tallec, Yann Ollivier. Unbiasing Truncated Backpropagation Through Time.

Updates of truncated BPTT

- ▶ **Principle:** Step k is represented by $(\theta^{(k)}, \tilde{h}_k)$. It is updated to $(\theta^{(k+1)}, \tilde{h}_{k+1})$ using the loss at time $k + 1$,

$$\begin{aligned}\theta^{(k+1)} &= \theta^{(k)} - \eta_k \tilde{\nabla}_{\theta} L_{k+1}(\theta^{(k)}), \\ \tilde{h}_{k+1} &= f(\tilde{h}_k, x_{k+1}; \theta^{(k)}),\end{aligned}$$

where **the truncated gradient**

$$\tilde{\nabla}_{\theta} L_{k+1}(\theta^{(k)}) = \nabla_{\theta} \ell(f(\tilde{h}_k, x_{k+1}; \theta), y_{k+1})|_{\theta=\theta_k}.$$

- ▶ **Difference to BPTT:** In BPTT, the step k is represented by $\theta^{(k)}$ and updated to $\theta^{(k+1)}$ using the loss over time $t \leq T$.

Example 4: BPTT vs. Truncated BPTT

- ▶ Consider the following problem:
 - ▶ $h_0 = x_0 \sim \mathcal{N}(\mu, \sigma^2)$
 - ▶ $h_t = f(h_{t-1}; \theta) = h_{t-1} + \theta$
 - ▶ $L_t = \ell(h_t) = \mathbb{E}_{x_0 \sim p}(h_t - \mu)^2$
 - ▶ $\min_{\theta \in \mathbb{R}} L = \sum_{t=1}^T L_t$
- ▶ What is the optimal solution of L ?
- ▶ What is the gradient $\nabla_{\theta} L$ and the truncated gradient $\tilde{\nabla}_{\theta} L_t$?